

1. Handling large data structures

During the classes we were trying to find the limit of the length of the vector, that we could have added.

At the beginning we were slowly increasing the length from 50 up to 1 000 000 000. The length of 10^9 was too large for our GPU. During the test we noticed, that the memory of our GPU we were using when we were adding vectors of length of 10^5 , was around 1GB. Knowing that our GPU has 6GB of memory, we were able to add vectors of 500 millions elements.

In the next phase of the classes, basing on our device properties (size of the memory - 6 371 475 456 bytes), and the size of the float (4 bytes), and the knowledge, that we have to find memory, for three vectors, we calculated that 530 956 288 is the maximum length of our vector.

2. Memory management utility

Writing programs in CUDA requires thinking about memory management - allocating it in Device, Host and copying data between them due to the fact that they are physically different elements separated via PCIe.

However with usage of CUDA Unified Memory, instead of having two different pools of memory - accessed separately by System and Device, we create single memory address space which is shared between both CPU and GPU. Migration of data is done automatically when it is needed. What is worth noting, that the problem with allocating and copying is not gone entirely, it is just hidden and performed by the utility. With Pascal architecture Unified Memory has hardware support for virtual memory page faulting and migration with Page Migration engine.

The main reason why do we use Unified Memory is simplicity. As almost the whole work regarding allocating data is done by the system, we do not have to worry that something might have not been done in this case properly. Also the code becomes easier to read as there is less code which allocates and moves data.

Another benefit of using Memory Management is performance. Thanks to migrating data on demand we can achieve similar or sometimes even better levels of speed compared to using separate memory pools. This is all thanks to complex operations of CUDA driver and runtime.


We have conducted a simple test by comparing vectorAdd sample program execution time with and without using Unified Memory. Both programs were run with same size of array and same grid configuration. We conducted three tests on each version to decrease the probability of random events. Results are presented in table below:

Table 1: Execution time per function

	vectorAdd	cudaFree	cudaMalloc(Managed)	total
Default	1.617ms	1.093ms	128.36ms	175.27ms
With Unified Memory	38.198ms	6.7055ms	98.506ms	178.07ms

On the base of the tables above we created charts to visualize the data.

Execution time in [ms]



If we want to look at the comparison quality-wise, we can spot clear differences, between using Unified Memory, and the default functions.

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	51.33%	8.6400us	2	4.3200us	4.3200us	4.3200us	[CUDA memcpy HtoD]
	26.24%	4.4160us	1	4.4160us	4.4160us	4.4160us	vectorAdd(float const *, float const *, float*, int)
	22.43%	3.7760us	1	3.7760us	3.7760us	3.7760us	[CUDA memcpy DtoH]
API calls:	99.56%	265.78ms	3	88.593ms	6.9140us	265.76ms	cudaMalloc
	0.18%	487.56us	1	487.56us	487.56us	487.56us	cuDeviceTotalMem
	0.11%	297.59us	96	3.0990us	838ns	97.429us	cuDeviceGetAttribute
	0.05%	132.14us	3	44.046us	6.9840us	107.63us	cudaFree
	0.04%	119.22us	3	39.739us	21.092us	52.800us	cudaMemcpy
	0.02%	65.931us	1	65.931us	65.931us	65.931us	cudaLaunchKernel
	0.02%	42.743us	1	42.743us	42.743us	42.743us	cuDeviceGetName
	0.00%	10.336us	1	10.336us	10.336us	10.336us	cuDeviceGetPCIBusId
	0.00%	9.4290us	1	9.4290us	9.4290us	9.4290us	cuDeviceSynchronize
	0.00%	4.1900us	3	1.3960us	1.0470us	2.0950us	cuDeviceGetCount
	0.00%	2.6540us	1	2.6540us	2.6540us	2.6540us	cudaGetLastError
	0.00%	2.5840us	2	1.2920us	1.1170us	1.4670us	cuDeviceGet
	0.00%	1.2570us	1	1.2570us	1.2570us	1.2570us	cuDeviceGetUuid

Profiling results - Default (not connected with the chart above)

Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	100.00%	955.66us	1	955.66us	955.66us	955.66us	vectorAdd(float const *, float const *, float*, int)
API calls:	99.27%	310.01ms	3	103.34ms	9.1490us	309.95ms	cudaMallocManaged
	0.31%	956.13us	1	956.13us	956.13us	956.13us	cuDeviceSynchronize
	0.15%	468.99us	1	468.99us	468.99us	468.99us	cuDeviceTotalMem
	0.09%	290.47us	96	3.0250us	838ns	94.146us	cuDeviceGetAttribute
	0.07%	203.80us	1	203.80us	203.80us	203.80us	cuDeviceGetProperties
	0.05%	153.02us	3	51.007us	17.810us	88.070us	cudaFree
	0.04%	115.94us	1	115.94us	115.94us	115.94us	cudaLaunchKernel
	0.01%	39.181us	1	39.181us	39.181us	39.181us	cuDeviceGetName
	0.01%	28.845us	1	28.845us	28.845us	28.845us	cudaSetDevice
	0.00%	10.266us	1	10.266us	10.266us	10.266us	cuDeviceGetPCIBusId
	0.00%	4.4700us	1	4.4700us	4.4700us	4.4700us	cuDeviceGetCount
	0.00%	3.8410us	3	1.2800us	978ns	1.8850us	cuDeviceGetCount
	0.00%	3.4220us	1	3.4220us	3.4220us	3.4220us	cudaGetLastError
	0.00%	2.5140us	2	1.2570us	978ns	1.5360us	cuDeviceGet
	0.00%	1.1180us	1	1.1180us	1.1180us	1.1180us	cuDeviceGetUuid

==40695== Unified Memory profiling result:						
Device "GeForce GTX 1060 6GB (0)"						
Count	Avg Size	Min Size	Max Size	Total Size	Total Time	Name
4	32.000KB	12.000KB	52.000KB	128.000KB	13.92000us	Host To Device
5	25.600KB	4.000KB	60.000KB	128.000KB	12.80000us	Device To Host

Profiling results - Unified Memory (not connected with the chart above)

Processes distinctive for each way:

Unified Memory:

`cudaMallocManaged`

Default:

`cudaMemcpy`

`cudaMalloc`

CUDA `memcpy DtoH`

CUDA `memcpy HtoD`

The difference is obvious after the explanation in the beginning of the second section.

Implementation of Unified Memory is fairly simple. Instead of using `malloc()` to allocate data, We use `cudaMallocManaged(void** ptr, size_t size)`, a function which returns pointer `ptr` which can be accessed by device or host. Since the memory is shared, we do not need to create different vectors for device and host. Furthermore we do not need to copy data, it means that any `cudaMemcpy` functions are no longer required. Using this method sparks another problem - synchronization. Accessing unprocessed data by host might be a major source of problems. Function `cudaDeviceSynchronize()` allows compute device to finish requested tasks by blocking host thread. After all the work is done we still need to free memory by using `cudaFree()`.

In the table below we have shown differences between both approaches of managing memory:

Table 2: Differences between code without and with Unified Memory.

Default	With Unified Memory
<code>float *h_A = (float *)malloc(size);</code>	<code>float *A;</code> <code>cudaMallocManaged(&A, numElements*sizeof(float));</code>
<code>cudaMalloc((void **) &d_A, size);</code>	--
--	<code>cudaDeviceSynchronize();</code>
<code>cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);</code>	--
<code>cudaMemcpy(h_A, d_A, size, cudaMemcpyDeviceToHost);</code>	--
<code>cudaFree(d_A);</code>	<code>cudaFree(A);</code>

Profiling program which uses Memory Management provides us with slightly more information. In the table there is a new section, where we can see the result of Unified Memory profiling.

We conducted a small test to get to know it a bit better.

Using previously mentioned programs, we added vectors, which had 10, 10^3 , 10^5 , and 10^7 elements. The profiling results are shown below.

```
==12591== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
  Count Avg Size Min Size Max Size Total Size Total Time Name
    2 32.000KB 4.0000KB 60.000KB 64.00000KB 6.912000us Host To Device
    2 32.000KB 4.0000KB 60.000KB 64.00000KB 6.112000us Device To Host
    1 - - - - 528.6080us Gpu page fault groups
Total CPU Page faults: 2
==12902== NVPROF is profiling process 12902, command: ./Lab3-01
```

10-element vector

```
==12902== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
  Count Avg Size Min Size Max Size Total Size Total Time Name
    2 32.000KB 4.0000KB 60.000KB 64.00000KB 7.040000us Host To Device
    2 32.000KB 4.0000KB 60.000KB 64.00000KB 6.112000us Device To Host
    1 - - - - 543.3600us Gpu page fault groups
Total CPU Page faults: 2
==13085== NVPROF is profiling process 13085, command: ./Lab3-01
```

1 000-element vector

```
==13085== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
  Count Avg Size Min Size Max Size Total Size Total Time Name
   24 42.666KB 4.0000KB 244.00KB 1.000000MB 101.1200us Host To Device
   21 97.523KB 4.0000KB 0.9961MB 2.000000MB 169.7920us Device To Host
    7 - - - - 1.365952ms Gpu page fault groups
Total CPU Page faults: 16
==13257== NVPROF is profiling process 13257, command: ./Lab3-01
```

100 000-element vector

```
==13257== Unified Memory profiling result:
Device "GeForce GTX 1060 6GB (0)"
  Count Avg Size Min Size Max Size Total Size Total Time Name
   625 125.00KB 4.0000KB 0.9883MB 76.29688MB 7.336480ms Host To Device
   696 168.38KB 4.0000KB 0.9961MB 114.4453MB 9.628256ms Device To Host
   317 - - - - 36.50758ms Gpu page fault groups
Total CPU Page faults: 580
```

10 000 000-element vector

We can clearly see that for smaller amounts of data (10, 10^3 elements), Unified Memory utility uses the same amount of memory and virtually the same amount of time to transfer data.

However for larger vectors, we can see almost proportional growth of memory usage (between 10^5 and 10^7 -element vectors), and significant increase in time consumption. Next to the “total size”, and “total time”, we can see the “min size”, and “max size” columns. We assume that these refer to the sizes of blocks of memory our utility uses, thus we can see not directly proportional growth between 10^5 and 10^7 vectors in memory consumption.

3.How to prevent processing too large amount of data

Running GPU programs require knowledge about the hardware you are running on: the amount of threads and memory you have available. Without it we may try to do computing which is physically impossible. This might happen f.eg. when we exceed total available memory.

In our vectorAdd code, trying to add vectors which in total have over 6GB of global memory will result in a crash, the program will simply stop responding. As programmers we know our limitations and what will happen when we fetch too much data, so naturally we avoid them. But the end user might not. In this case we have to introduce an algorithm which will prevent overloading the memory.

Firstly we gather information about hardware we possess. With DeviceQuery we get the number of GPU's and how much memory do they have. In this example, knowledge that we have an addition of 1D arrays $A+B=C$, allows us to calculate how much elements each of them can consist of. Dividing total available memory by the number of vectors and size of data type (rounded down) results in the maximal size of arrays.

With that we create an if statement which checks the number of elements given by user. If it's larger than calculated limit, there are two ways of dealing with this problem. First one: we show message that memory will be overloaded and exit the program with `exit(EXIT_FAILURE)`. Another one is where we cap the vectors with maximal possible size, ignore the rest of data and print appropriate statement.

The solution we have chosen is presented in the fragment of code below:

```
float totalmem=0;
int deviceCount = 0;
int dev;
cudaGetDeviceCount(&deviceCount);
for (dev = 0; dev < deviceCount; ++dev) {
    cudaSetDevice(dev);
    cudaDeviceProp deviceProp;
    cudaGetDeviceProperties(&deviceProp, dev);
    totalmem+=deviceProp.totalGlobalMem;
}

float memperarr=(int) (deviceCount*totalmem/12);
//...
if(numElements>memperarr){
    printf("Size of array exceeds total device memory!\nMaximal number of
elements in array is %f",memperarr);
    printf("\nArray dimension will be changed to prevent crash\nWARNING! %f
elements will be lost!!\n\n",numElements-memperarr);
    numElements=memperarr;
```

